




CSE446: Instance-based Learning (a.k.a. non-parametric methods) Winter 2015

Luke Zettlemoyer

Slides adapted from Carlos Guestrin


Linear Regression: What can go wrong?




What do we do if the bias is too strong?

- Might want the data to drive the complexity of the model!
- Try instance-based Learning (a.k.a. non-parametric methods)?

Using data to predict new data



Nearest neighbor with lots of data!



Univariate 1-Nearest Neighbor

Given data $(x^1, y^1) (x^2, y^2) .. (x^N, y^N)$, where we assume $y=f(x)$ for some unknown function f .

Given query point x , your job is to predict $y=f(x)$


Nearest Neighbor:

1. Find the closest x^i in our set of datapoints

$$i(nn) = \underset{i}{\operatorname{argmin}} |x^i - x|$$

2. Predict $y^{i(nn)}$

Here's a dataset with one input, one output and four datapoints.




1-Nearest Neighbor is an example of....

Instance-based learning

A function approximator that has been around since about 1910.

To make a prediction, search database for similar datapoints, and fit with the local points.



Instance-based learning, four things to specify:

- A distance metric
- How many nearby neighbors to look at?
- A weighting function (optional)
- How to fit with the local points?

1-Nearest Neighbor

Instance-based learning, four things to specify:

1. A distance metric

Often Euclidian (many more are possible)

2. How many nearby neighbors to look at?

One

3. A weighting function (optional)


Unused

4. How to fit with the local points?


Just predict the same output as the nearest neighbor.

Multivariate 1-NN examples


Classification




Regression




Notable distance metrics (and their level sets)




Weighted Euclidian (L_2)



L_1 norm (absolute)



Mahalanobis



L_∞ (max) norm







Consistency of 1-NN

- Consider an estimator f_n trained on n examples
 - e.g., 1-NN, neural nets, regression,...
- Estimator is *consistent* if true error goes to zero as amount of data increases
 - e.g., for no noise data, consistent if for any data distribution $p(x)$:

$$\lim_{n \rightarrow \infty} MSE(f_n) = 0 \quad MSE(f_n) = \int_x p(x) (f_n(x) - y_x)^2 dx$$

- Linear regression is not consistent!
 - Representation bias
- **1-NN is consistent**
 - What about noisy data?
 - What about variance?

1-NN overfits?



k-Nearest Neighbor

Instance-based learning, four things to specify:

1. *A distance metric*

Euclidian (and many more)

2. *How many nearby neighbors to look at?*

k

1. *A weighting function (optional)*

Unused




2. *How to fit with the local points?*

Return the average output




predict: $(1/k) \sum_i y^i$ (summing over k nearest neighbors)

k-Nearest Neighbor

k=1



k=9





Which is better? What can we do about the discontinuities?

Weighted distance metrics

Suppose the input vectors $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^N$ are two dimensional:

$$\mathbf{x}^1 = (x^1_1, x^1_2), \mathbf{x}^2 = (x^2_1, x^2_2), \dots, \mathbf{x}^N = (x^N_1, x^N_2).$$

Nearest-neighbor regions in input space:



$$Dist(\mathbf{x}^i, \mathbf{x}^j) = (x^i_1 - x^j_1)^2 + (x^i_2 - x^j_2)^2$$

$$Dist(\mathbf{x}^i, \mathbf{x}^j) = (x^i_1 - x^j_1)^2 + (3x^i_2 - 3x^j_2)^2$$

The relative scaling of the distance metric affect region shapes

Weighted Euclidean distance metric

Or equivalently,

$$D(\mathbf{x}, \mathbf{x}') = \sqrt{\sum_i \sigma_i^2 (x_i - x'_i)^2}$$

where

$$D(\mathbf{x}, \mathbf{x}') = \sqrt{(\mathbf{x} - \mathbf{x}')^T \Sigma (\mathbf{x} - \mathbf{x}')}$$

$$\Sigma = \begin{bmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & 0 \\ \vdots & \vdots & \dots & \vdots \\ 0 & 0 & \dots & \sigma_N^2 \end{bmatrix}$$

Other Metrics...

- Mahalanobis, Rank-based, Correlation-based,...

Kernel regression

Instance-based learning:

1. *A distance metric*

Euclidian (and many more)

2. *How many nearby neighbors to look at?*

All of them

3. *A weighting function*


$$w^i = \exp(-D(x^i, \text{query})^2 / K_w^2)$$

Nearby points to the query are weighted strongly, far points weakly. The K_w parameter is the **Kernel Width**. Very important.

4. *How to fit with the local points?*

Predict the weighted average of the outputs:

$$\text{predict} = \Sigma w^i y^i / \Sigma w^i$$

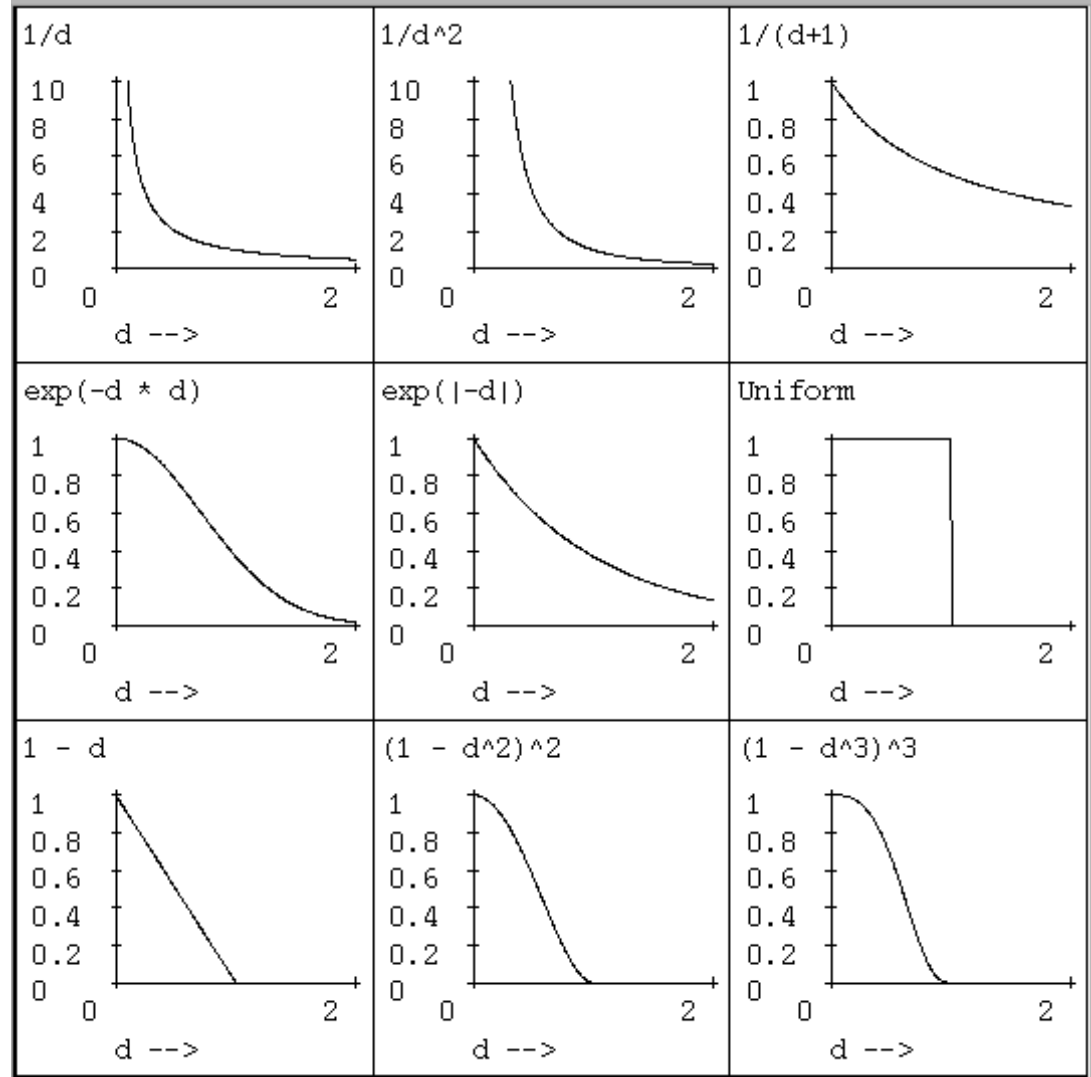


Many possible weighting functions

$$w^i = \exp(-D(x^i, \text{query})^2 / K_w^2)$$


Typically:

- Choose D manually
- Optimize K_w using gradient descent




(Our examples use Gaussian)


Kernel regression predictions



$K_W=10$



$K_W=20$






$K_W=80$

Increasing the kernel width K_W means further away points get an opportunity to influence you.




As $K_W \rightarrow \infty$, the prediction tends to the global average.

Kernel regression on our test cases

NN k=9



Kernel regression




$K_W = 1/32$ of x-axis width.

$K_W = 1/32$ of x-axis width.


$K_W = 1/16$ axis width.

Choosing a good K_W is important! Remind you of anything we have seen?


Kernel regression: problem solved?



K_W = Best.



K_W = Best.



K_W = Best.

Where are we having problems?

- Sometimes in the middle...
- Generally, on the ends (extrapolation is hard!)

Time to try something more powerful...!!!

Locally weighted regression

Kernel regression:

- Take a very very conservative function approximator called AVERAGING.
- Locally weight it.

Locally weighted regression:

- Take a conservative function approximator called LINEAR REGRESSION.
- Locally weight it.

Locally weighted regression


Instance-based learning, four things to specify:

- *A distance metric*
Any
- *How many nearby neighbors to look at?*
All of them
- *A weighting function (optional)*
Kernels: $w^i = \exp(-D(x_i, \text{query})^2 / Kw^2)$
- *How to fit with the local points?*

General weighted regression:

$$\hat{w} = \underset{w}{\operatorname{argmin}} \sum_{k=1}^N (w^k (y^k - w^T x^k))^2$$


How LWR works



Linear regression

- Same parameters for all queries

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$



Solving for each input: complex surface!


Locally weighted regression

- Solve weighted linear regression for each query


$$\beta = \left((WX)^T WX \right)^{-1} (WX)^T WY$$
$$W = \begin{pmatrix} w_1 & 0 & 0 & 0 \\ 0 & w_2 & 0 & 0 \\ 0 & 0 & \ddots & 0 \\ 0 & 0 & 0 & w_n \end{pmatrix}$$

LWR on our test cases


Kernel regression



$K_W = 1/32$ of x-axis width.




$K_W = 1/32$ of x-axis width.




$K_W = 1/16$ axis width.


LWR



$K_W = 1/16$ of x-axis width.




$K_W = 1/32$ of x-axis width.




$K_W = 1/8$ of x-axis width.

Locally weighted polynomial regression


Kernel Regression: Kernel width K_W at optimal level.



$K_W = 1/100$ x-axis



$K_W = 1/40$ x-axis




$K_W = 1/15$ x-axis

Local quadratic regression is easy: just add quadratic terms to the WXTWX matrix. As the regression degree increases, the kernel width can increase without introducing bias.

Challenges for instance based learning

- Must store and retrieve all data!
 - Most real work done during testing
 - For every test sample, must search through all dataset – very slow!
 - But, there are fast methods for dealing with large datasets
- Instance-based learning often poor with noisy or irrelevant features
 - In high dimensional spaces, all points will be very far from each other
 - Typically need a number of examples that is exponential in the dimension of X
 - But, sometimes you are ok if you are clever about features

Curse of the irrelevant feature



This is a contrived example, but similar problems are common in practice
Need some form of feature selection!!

What you need to know about instance-based learning

- **k-NN**
 - Simplest learning algorithm
 - With sufficient data, very hard to beat “strawman” approach
 - Picking k ?
- **Kernel regression**
 - Set k to n (number of data points) and optimize weights by gradient descent
 - Smoother than k-NN
- **Locally weighted regression**
 - Generalizes kernel regression, not just local average
- **Curse of dimensionality**
 - Must remember (very large) dataset for prediction
 - Irrelevant features often killers for instance-based approaches

Acknowledgment

- This lecture contains some material from Andrew Moore's excellent collection of ML tutorials:
 - <http://www.cs.cmu.edu/~awm/tutorials>