

CSE446: PAC-learning, VC Dimension Winter 2015

Luke Zettlemoyer

Slides adapted from Carlos Guestrin

What now...

- We have explored *many* ways of learning from data
- But...
 - How good is our classifier, really?
 - How much data do I need to make it “good enough”?

A simple setting...


- Classification
 - m data points
 - **Finite** number of possible hypothesis (e.g., dec. trees of depth d)
- A learner finds a hypothesis h that is **consistent** with training data
 - Gets zero error in training – $error_{train}(h) = 0$
- What is the probability that h has more than ε true error?
 - $error_{true}(h) \geq \varepsilon$

How likely is a bad hypothesis to get m data points right?

- Hypothesis h that is **consistent** with training data
 - got m i.i.d. points right
 - h “bad” if it gets all this data right, but has high true error
 - What is the probability of this happening?
- Prob. h with $\text{error}_{\text{true}}(h) \geq \varepsilon$ gets randomly drawn data point right
$$P(\text{error}_{\text{true}}(h) \geq \varepsilon, \text{ gets one data point right}) \leq 1 - \varepsilon$$
- Prob. h with $\text{error}_{\text{true}}(h) \geq \varepsilon$ gets m iid data points right
$$P(\text{error}_{\text{true}}(h) \geq \varepsilon, \text{ gets } m \text{ iid data point right}) \leq (1 - \varepsilon)^m$$


But there are many possible hypothesis that are consistent with training data

- Which classifier should be learn?
 - and how to we generalize the bounds?
- We want to make as few assumptions as possible!
- So, pick any $h \in H_c$
- But wait, we had a bound on a single h , now we need to bound the worst $h \in H_c$



Union bound

- $P(A \text{ or } B \text{ or } C \text{ or } D \text{ or } \dots)$
 $\leq P(A) + P(B) + P(C) + P(D) + \dots$



Q: Is this a tight bound? Will it be useful?

How likely is learner to pick a bad hypothesis

$$P(\text{error}_{\text{true}}(h) \geq \varepsilon, \text{ gets } m \text{ iid data point right}) \leq (1-\varepsilon)^m$$

There are k hypothesis consistent with data

- How likely is learner to pick a bad one?
- We need to a bound that holds for all of them!

$$P(\text{error}_{\text{true}}(h_1) \geq \varepsilon \text{ OR } \text{error}_{\text{true}}(h_1) \geq \varepsilon \text{ OR } \dots \text{ OR } \text{error}_{\text{true}}(h_k) \geq \varepsilon)$$

$$\begin{aligned} &\leq \sum_k P(\text{error}_{\text{true}}(h_k) \geq \varepsilon) && \leftarrow \text{Union bound} \\ &\leq \sum_k (1-\varepsilon)^m && \leftarrow \text{bound on individual } h_j\text{s} \\ &\leq |H|(1-\varepsilon)^m && \leftarrow k \leq |H| \\ &\leq |H| e^{-m\varepsilon} && \leftarrow (1-\varepsilon) \leq e^{-\varepsilon} \text{ for } 0 \leq \varepsilon \leq 1 \end{aligned}$$

Generalization error in finite hypothesis spaces [Haussler '88]

- **Theorem:** Hypothesis space H finite, dataset D with m i.i.d. samples, $0 < \epsilon < 1$: for any learned hypothesis h that is consistent on the training data:

$$P(\text{error}_{\text{true}}(h) > \epsilon) \leq |H|e^{-m\epsilon}$$

Using a PAC bound

- Typically, 2 use cases:
 - 1: Pick ϵ and δ , compute m
 - 2: Pick m and δ , compute ϵ

Argument: For all h we know that

$$P(\text{error}_{\text{true}}(h) > \epsilon) \leq |H|e^{-m\epsilon}$$

so, with probability $1-\delta$ the following holds...

$$P(\text{error}_{\text{true}}(h) \leq \epsilon) \leq |H|e^{-m\epsilon} \leq \delta$$

$$\ln(|H|e^{-m\epsilon}) \leq \ln \delta$$

Case 1 $\ln |H| - m\epsilon \leq \ln \delta$

$$m \geq \frac{\ln |H| + \ln \frac{1}{\delta}}{\epsilon}$$

Log dependence on $|H|$,
ok if exponential size (but
not doubly)

ϵ has stronger
influence than δ

Case 2

$$\epsilon \geq \frac{\ln |H| + \ln \frac{1}{\delta}}{m}$$

ϵ shrinks at rate $O(1/m)$

Limitations of Haussler '88 bound

$$P(\text{error}_{\text{true}}(h) > \epsilon) \leq |H|e^{-m\epsilon}$$

- Do we really want to pick a consistent hypothesis h ? (where $\text{error}_{\text{train}}(h)=0$)
- Size of hypothesis space
 - What if $|H|$ is really big?
 - What if it is continuous?
- **First Goal:** Can we get a bound for a learner with $\text{error}_{\text{train}}(h)$ in training set?

Question: What's the expected error of a hypothesis?

- The error of a hypothesis is like estimating the parameter of a coin!
- Chernoff bound: for m i.i.d. coin flips, x_1, \dots, x_m , where $x_i \in \{0, 1\}$. For $0 < \epsilon < 1$:

$$P \left(\theta - \frac{1}{m} \sum_i x_i > \epsilon \right) \leq e^{-2m\epsilon^2}$$

Generalization bound for $|H|$ hypothesis

- **Theorem:** Hypothesis space H finite, dataset D with m i.i.d. samples, $0 < \epsilon < 1$: for any learned hypothesis h :

$$P(\text{error}_{\text{true}}(h) - \text{error}_{\text{train}}(h) > \epsilon) \leq |H|e^{-2m\epsilon^2}$$

Why? Same reasoning as before. Use the Union bound over individual Chernoff bounds

PAC bound and Bias-Variance tradeoff

$$P(\text{error}_{\text{true}}(h) - \text{error}_{\text{train}}(h) > \epsilon) \leq |H|e^{-2m\epsilon^2}$$

**or, after moving some terms around,
with probability at least $1-\delta$:**

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}$$

Important: PAC bound holds for all h , but doesn't guarantee that algorithm finds best h !!!

PAC bound and Bias-Variance tradeoff

for all h , with probability at least $1-\delta$:

$$\text{error}_{\text{true}}(h) \leq \underbrace{\text{error}_{\text{train}}(h)}_{\text{"bias"}} + \underbrace{\sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}}_{\text{"variance"}}$$

- For large $|H|$
 - low bias (assuming we can find a good h)
 - high variance (because bound is looser)
- For small $|H|$
 - high bias (is there a good h ?)
 - low variance (tighter bound)

PAC bound: How much data?

$$P(\text{error}_{\text{true}}(h) - \text{error}_{\text{train}}(h) > \epsilon) \leq |H|e^{-2m\epsilon^2}$$

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}$$

- Given δ, ϵ how big should m be?

$$m \geq \frac{1}{2\epsilon^2} \left(\ln |H| + \ln \frac{1}{\delta} \right)$$

Decision Trees

$$m \geq \frac{1}{2\epsilon^2} \left(\ln |H| + \ln \frac{1}{\delta} \right)$$

- Bound number of decision trees with depth k with data that has n features:

$$2 * (2n)^{2^k - 1}$$

- **Bad!!!** Need exponentially many data points (in k)!!!

$$m \geq \frac{\ln 2}{2\epsilon^2} \left((2^k - 1)(1 + \log_2 n) + 1 + \ln \frac{1}{\delta} \right)$$

- But, for m data points, tree can't get too big...
 - Number of leaves never more than number data points
 - Instead, let's bound number of decision trees with k leaves

$$H_k = n^{k-1} (k + 1)^{2^k - 1}$$

PAC bound for decision trees with k leaves – Bias-Variance revisited

$$H_k = n^{k-1} (k+1)^{2k-1} \quad \text{error}_{true}(h) \leq \text{error}_{train}(h) + \sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}$$

$$\text{error}_{true}(h) \leq \text{error}_{train}(h) + \sqrt{\frac{(k-1) \ln n + (2k-1) \ln(k+1) + \ln \frac{1}{\delta}}{2m}}$$

Bias / variance again

- $k \ll m$: high bias, low variance
- $k=m$: no bias, high variance
- $k>m$: we would never do this!!!

What did we learn from decision trees?

- Bias-Variance tradeoff formalized

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{(k-1) \ln n + (2k-1) \ln(k+1) + \ln \frac{1}{\delta}}{2m}}$$

- Moral of the story:

Complexity of learning not measured in terms of size hypothesis space, but in maximum *number of points* that allows consistent classification


What about continuous hypothesis spaces?

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{\ln |H| + \ln \frac{1}{\delta}}{2m}}$$


- Continuous hypothesis space:
 - $|H| = \infty$
 - Infinite variance???
- **As with decision trees, only care about the maximum number of points that can be classified exactly!**

How many points can a linear boundary classify exactly? (1-D)

2 Points: Yes!!



3 Points: No...



etc (8 total)

Shattering and VC Dimension

A set of points is *shattered* by a hypothesis space H iff:


- For all ways of *splitting* the examples into positive and negative subsets
- There exists some *consistent* hypothesis h

The *VC Dimension* of H over input space X


- The size of the *largest* finite subset of X shattered by H

How many points can a linear boundary classify exactly? (2-D)

3 Points: Yes!!



4 Points: No...



How many points can a linear boundary classify exactly? (d-D)

- A linear classifier $w_0 + \sum_{j=1..d} w_j x_j$ can represent all assignments of possible labels to $d+1$ points
 - But not $d+2$!!
 - Bias term w_0 required!
 - **Rule of Thumb:** number of parameters in model often matches max number of points
- **Question:** Can we get a bound for error in as a function of the number of points that can be completely labeled?

PAC bound using VC dimension

- **VC dimension:** number of training points that can be classified exactly (shattered) by hypothesis space H !!!
 - Measures relevant size of hypothesis space, as with decision trees with k leaves

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{VC(H) \left(\ln \frac{2m}{VC(H)} + 1 \right) + \ln \frac{4}{\delta}}{m}}$$

- **Same bias / variance tradeoff as always**
 - Now, just a function of $VC(H)$

Examples of VC dimension

$$\text{error}_{\text{true}}(h) \leq \text{error}_{\text{train}}(h) + \sqrt{\frac{VC(H) \left(\ln \frac{2m}{VC(H)} + 1 \right) + \ln \frac{4}{\delta}}{m}}$$

- Linear classifiers:
 - $VC(H) = d+1$, for d features plus constant term b
- Neural networks (we will see this next)
 - $VC(H) = \# \text{parameters}$
 - Local minima means NNs will probably not find best parameters
- 1-Nearest neighbor
 - $VC(H) = \infty$
- SVM with Gaussian Kernel
 - $VC(H) = \infty$

What you need to know

- Finite hypothesis space
 - Derive results
 - Counting number of hypothesis
 - Mistakes on Training data
- Complexity of the classifier depends on number of points that can be classified exactly
 - Finite case – decision trees
 - Infinite case – VC dimension
- Bias-Variance tradeoff in learning theory
- Remember: will your algorithm find best classifier?