

CSE 446 Machine Learning, Winter 2015

Homework 4

Due: March 13, 2015

1 Principal Component Analysis by hand [15 points]

You are given the following 3 data points in 2-d space.

- $A = (1,1)$
- $B = (4,5)$
- $C = (7,9)$

Please show your work for all of the questions below.

1. [5 points] What is the first principal Eigenvector of the data points?
2. [5 points] If we project the data points onto the 1-d space by PCA, what would be the variance of projected data?
3. [5 points] What is the reconstruction error if we represent the projected points on the original 2-d space? Briefly explain your answer.

2 Manual calculation of one round of EM for a GMM [40 points]

(Extended version of: Murphy Exercise 11.7) In this question we consider clustering 1D data with a mixture of 2 Gaussians using the EM algorithm. You are given the 1-D data points $x = [1 \ 10 \ 20]$.

M step

Suppose the output of the E step is the following matrix:

$$R = \begin{pmatrix} 1 & 0 \\ 0.4 & 0.6 \\ 0 & 1 \end{pmatrix}$$

where entry $R_{i,c}$ is the probability of observation x_i belonging to cluster c (the responsibility of cluster c for data point i). You just have to compute the M step. You may state the

equations for maximum likelihood estimates of these quantities (which you should know) without proof; you just have to apply the equations to this data set. You may leave your answer in fractional form. Show your work.

1. [4 points] Write down the likelihood function you are trying to optimize.
2. [8 points] After performing the M step for the mixing weights π_1, π_2 , what are the new values?
3. [8 points] After performing the M step for the means μ_1 and μ_2 , what are the new values?
4. [8 points] After performing the M step for the standard deviations σ_1 and σ_2 , what are the new values?

E step

Now suppose the output of the M step is the answer to the previous section. You will compute the subsequent E step.

1. [4 points] Write down the formula for the probability of observation x_i belonging to cluster c .
2. [8 points] After performing the E step, what is the new value of R ?

3 Programing Question (K-means) [45 points]

3.1 The Data

There are two files with the data. The first `digit.txt` contains the 1000 observations of 157 pixels (a subset of the original 785) from images containing handwritten digits. The second file `labels.txt` contains the true digit label (either 1, 3, 5, or 7). You can read both data files in with

```
X = genfromtxt('digit.txt')
Y = genfromtxt('labels.txt', dtype=int)
```

Please note that there aren't IDs for the digits. Please assume the first line is ID 0, the second line is ID 1, and so on. The labels correspond to the digit file, so the first line of `labels.txt` is the label for the digit in the first line of `digit.txt`.

3.2 The algorithm

Here is a quick review of how K-means algorithm works.

1. Select k starting centers that are points from your data set. You should be able to select these centers randomly or have them given as a parameter.
2. Assign each data point to the cluster associated with the nearest of the k center points.
3. Re-calculate the centers as the mean vector of each cluster from (2).

4. Repeat steps (2) and (3) until convergence or iteration limit.

Define convergence as no change in label assignment from one step to another **or** you have iterated 20 times (whichever comes first). Please count your iterations as follows: after 20 iterations, you should have assigned the points 20 times.

For Step 1, please make sure that you deep-copy your data points when you select them as starting centers. Otherwise, when you move centers, you will end up moving the data points as well.

3.3 Within group sum of squares

The goal of clustering can be thought of as minimizing the variation within groups and consequently maximizing the variation between groups. A good model has low sum of squares within each group.

We define sum of squares in the traditional way. Let C_k be the k th cluster and let μ_k be the empirical mean of the observations x_i in cluster C_k . Then the within group sum of squares for cluster C_k is defined as:

$$SS(k) = \sum_{i \in C_k} |x_i - \mu_{C_k}|^2$$

Please note that the term $|x_i - \mu_{C_k}|$ is the euclidean distance between x_i and μ_{C_k} , and therefore should be calculated as $|x_i - \mu_{C_k}| = \sqrt{\sum_{j=1}^d (x_{ij} - \mu_{C_{kj}})^2}$, where d is the number of dimensions. Please note that that term is squared in $SS(k)$.

If there are K clusters total then the “sum of within group sum of squares” is just the sum of all K of these individual $SS(k)$ terms.

3.4 Mistake Rate

Given that we know the actual assignment labels for each data point we can attempt to analyze how well the clustering recovered this. For cluster C_k let its assignment be whatever the majority vote is for that cluster. If there is a tie, just choose the digit that is smaller numerically as the majority vote.

For example if for one cluster we had 270 observations labeled **one**, 50 labeled **three**, 9 labeled **five**, and 0 labeled **seven** then that cluster will be assigned value **one** and had $50 + 9 + 0 = 59$ mistakes. If we add up the total number of “mistakes” for each cluster and divide by the total number of observations, we will get our total mistake rate, between 0 and 1.

3.5 Questions

When you have implemented the algorithm please report the following:

1. [9pts] The values of sum of within group sum of squares and mistake rates for $k = 2$, $k = 4$ and $k = 6$. Please start your centers with the first k points in the dataset. So, if $k = 2$, your initial centroids will be ID 0 and ID 1, which correspond to the first two lines in the file.

2. [6pts] The number of iterations that k-means ran for $k = 6$, starting the centers as in the previous item. Make sure you count the iterations correctly. If you start with iteration $i = 0$ and at $i = 3$ the cluster assignments don't change, the number of iterations was 4, as you had to do step 2 four times to figure this out.
3. [15pts] A plot of the sum of within group sum of squares versus k for $k = 1, 2, 3, 4, 5, 6, 7, 8, 9, 10$. Please start your centers randomly (choose k points from the dataset at random).
4. [15pts] A plot of total mistake rate versus k for $k = 1, 2, 3, 4, 5, 6, 7, 8, 9, 10$. Please start your centers randomly (choose k points from the dataset at random).

For the last two items, you should generate these graphs a few times in case you get unlucky centers in the beginning, but please submit only one plot.